



# MIE524 Data Mining

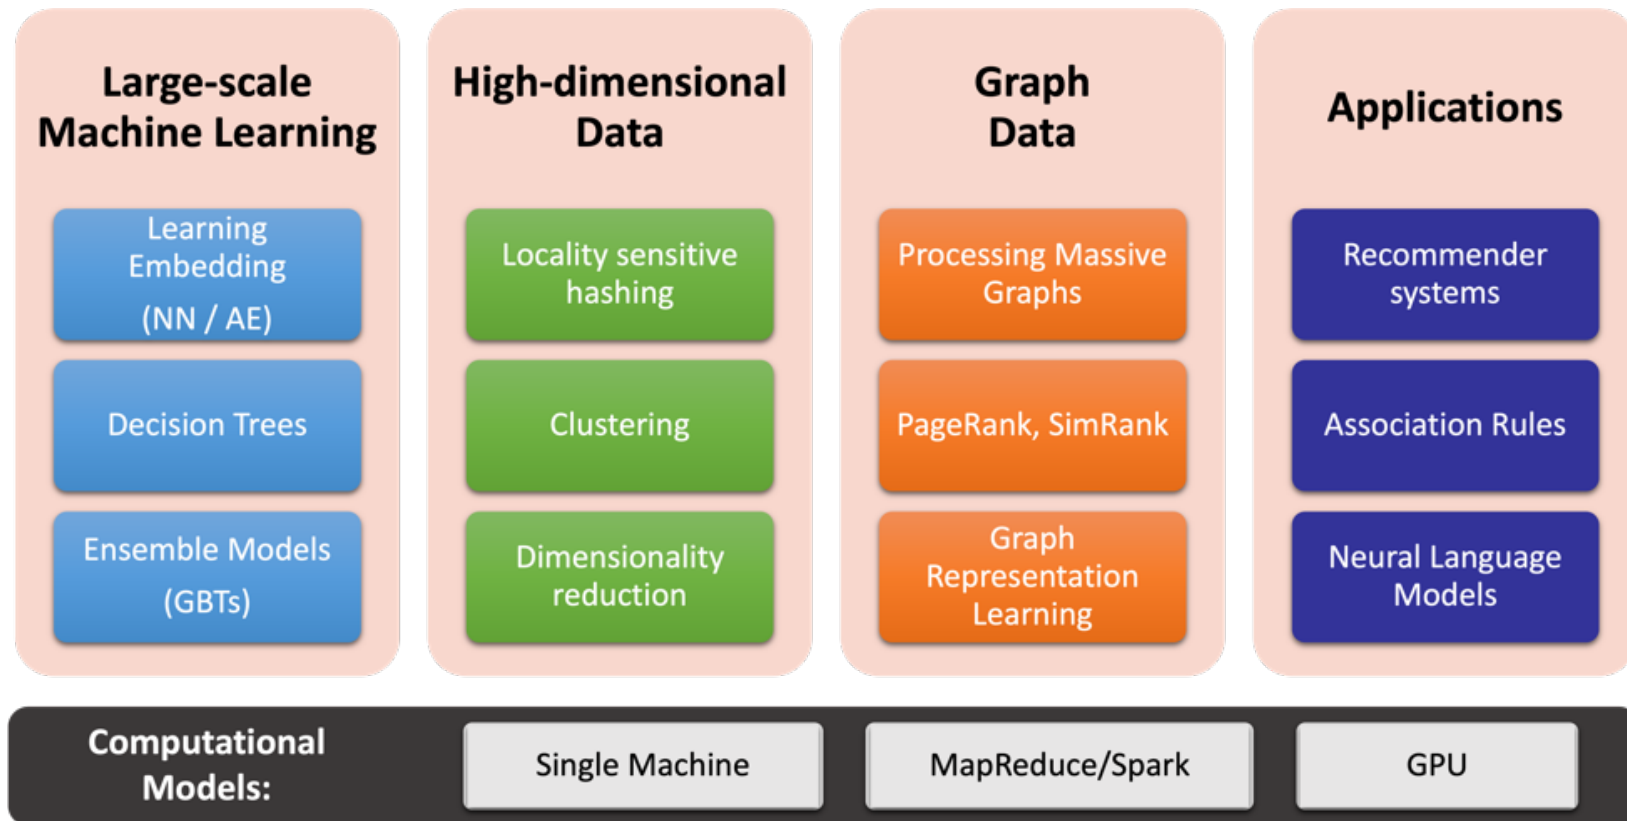
## **Clustering**

Slides Credits:

Slides from Leskovec, Rajaraman, Ullman (<http://www.mmids.org>), Leskovec & Ghashami  
Alexandra Chouldechova, Roger Grosse, Amir-massoud Farahmand, and Juan Carrasquilla

# MIE524: Course Topics (Tentative)


---




# The Problem of Clustering

- Given a **set of points**, with a notion of **distance** between points, **group the points** into some number of *clusters*, so that
  - Members of the same cluster are close/similar to each other
  - Members of different clusters are dissimilar
- **Usually:**
  - Points are in a high-dimensional space
  - Similarity is defined using a distance measure
    - Euclidean, Cosine, Jaccard, edit distance, ...

# Example: Clusters & Outliers



# Example: Clusters & Outliers



## Examples of clustering tasks

- Identify similar groups of **online shoppers** based on their browsing and purchasing history
- Identify similar groups of **music listeners** or **movie viewers** based on their ratings or recent listening/viewing patterns
- Cluster **input variables** based on their correlations to remove redundant predictors from consideration
- Cluster hospital **patients** based on their medical histories
- Determine how to place **sensors**, **broadcasting towers**, **law enforcement**, or **emergency-care** centers to guarantee that desired *coverage criteria* are met

# Clustering Problem: Music CDs

- **Intuitively:** Music can be divided into categories, and customers prefer a few genres
  - But what are categories really?
- Represent a CD by a set of customers who bought it
- Similar CDs have similar sets of customers, and vice-versa

# Clustering Problem: Music CDs

## Space of all CDs:

- Think of a space with one dim. for each customer
  - Values in a dimension may be 0 or 1 only
  - A CD is a “point” in this space  $(x_1, x_2, \dots, x_d)$ , where  $x_i = 1$  iff the  $i^{\text{th}}$  customer bought the CD
- For Amazon, the dimension is tens of millions
- **Task:** Find clusters of similar CDs




# Why is it hard?

- Clustering in two dimensions looks easy
- Clustering small amounts of data looks easy
- And in most cases, looks are **not** deceiving
- Many applications involve not 2, but 10 or 10,000 dimensions
- **High-dimensional spaces look different:**  
Almost all pairs of points are very far from each other --> **The Curse of Dimensionality!**

# Example: Curse of Dimensionality

- Take 10,000 uniform random points on  $[0,1]$  line. Assume query point is at the origin
- What fraction of “space” do we need to cover to get 0.1% of data (10 nearest neighbors)
- In 1-dim to get 10 neighbors we must go to distance  $10/10,000=0.001$  on the average
- In 2-dim we must go  $\sqrt{0.001}=0.032$  to get a square that contains 0.001 volume
- In general, in  $d$ -dim we must go  $(0.001)^{\frac{1}{d}}$
- So, in 10-dim to capture 0.1% of the data we need 50% of the range.

# Curse of dimensionality





Formally, imagine the unit cube  $[0, 1]^d$ . All training data is sampled *uniformly* within this cube, i.e.  $\forall i, x_i \in [0, 1]^d$ , and we are considering the  $k = 10$  nearest neighbors of such a test point.


Let  $\ell$  be the edge length of the smallest hyper-cube that contains all  $k$ -nearest neighbor of a test point. Then  $\ell^d \approx \frac{k}{n}$  and  $\ell \approx \left(\frac{k}{n}\right)^{1/d}$ . If  $n = 1000$ , how big is  $\ell$ ?

$d$	$\ell$
2	0.1
10	0.63
100	0.955
1000	0.9954

# Example: Curse of Dimensionality

**Curse of Dimensionality:** All points are very far from each other





# *k*-means Clustering

## Method: K-mean clustering

- Main idea: A good clustering is one for which the within-cluster variation is as small as possible.
- The within-cluster variation for cluster  $C_k$  is some measure of the amount by which the observations within each class differ from one another
- We'll denote it by  $WCV(C_k)$
- Goal: Find  $C_1, \dots, C_K$  that minimize

$$\sum_{k=1}^K WCV(C_k)$$

- This says: Partition the observations into  $K$  clusters such that the WCV summed up over all  $K$  clusters is as small as possible

## How to define within-cluster variation?

- Goal: Find  $C_1, \dots, C_K$  that minimize

$$\sum_{k=1}^K \text{WCV}(C_k)$$

- Typically, we use squared Euclidean distance:

$$\text{WCV}(C_k) = \frac{1}{|C_k|} \sum_{i, i' \in C_k} \sum_{j=1}^p (x_{ij} - x_{i'j})^2$$


where  $|C_k|$  denotes the number of observations in cluster  $k$

- To be clear: We're treating  $K$  as fixed ahead of time. We are *not* optimizing  $K$  as part of this objective.

## Simple example

Here  $n = 5$  and  $K = 2$ ,  
The full *distance matrix* for all 5 observations is shown below.

	1	2	3	4	5
1	0	0.25	0.98	0.52	1.09
2	0.25	0	1.09	0.53	0.72
3	0.98	1.09	0	0.10	0.25
4	0.52	0.53	0.10	0	0.17
5	1.09	0.72	0.25	0.17	0



- **Red clustering:**  $\sum WCV_k = (0.25 + 0.53 + 0.52)/3 + 0.25/2 = 0.56$
- **Blue clustering:**  $\sum WCV_k = 0.25/2 + (0.10 + 0.17 + 0.25)/3 = 0.30$
- It's easy to see that the **Blue clustering** minimizes the within-cluster variation among all possible partitions of the data into  $K = 2$  clusters



## How do we minimize WCV?

$$\begin{aligned}\sum_{k=1}^K \text{WCV}(C_k) &= \sum_{k=1}^K \frac{1}{|C_k|} \sum_{i,i' \in C_k} \sum_{j=1}^p (x_{ij} - x_{i'j})^2 \\ &= \sum_{k=1}^K \frac{1}{|C_k|} \sum_{i,i' \in C_k} \|x_i - x_{i'}\|_2^2\end{aligned}$$

- It's *computationally infeasible* to actually minimize this criterion
- We essentially have to try all possible partitions of  $n$  points into  $K$  sets.
- When  $n = 10$ ,  $K = 4$ , there are 34,105 possible partitions
- When  $n = 25$ ,  $K = 4$ , there are  $5 \times 10^{13}$ ...
- We're going to have to settle for an **approximate solution**

# K-means algorithm

- It turns out that we can rewrite  $WCV_k$  more conveniently:

$$WCV_k = \frac{1}{|C_k|} \sum_{i,i' \in C_k} \|x_i - x_{i'}\|_2^2 = 2 \sum_{i \in C_k} \|x_i - \bar{x}_k\|^2$$

where  $\bar{x}_k = \frac{1}{|C_k|} \sum_{i \in C_k} x_i$  is just the average of all the points in cluster  $C_k$

# $k$ -means Algorithm(s)

- Initialize clusters by picking  $k$  centers

## Until convergence:

- **1)** For each point, assign it to the cluster whose current centroid is the closest
- **2)** After all points are assigned, update the centroids of the  $k$  clusters as **average of datapoints within each cluster**

**Convergence means** Points don't move between clusters and centroids stabilize

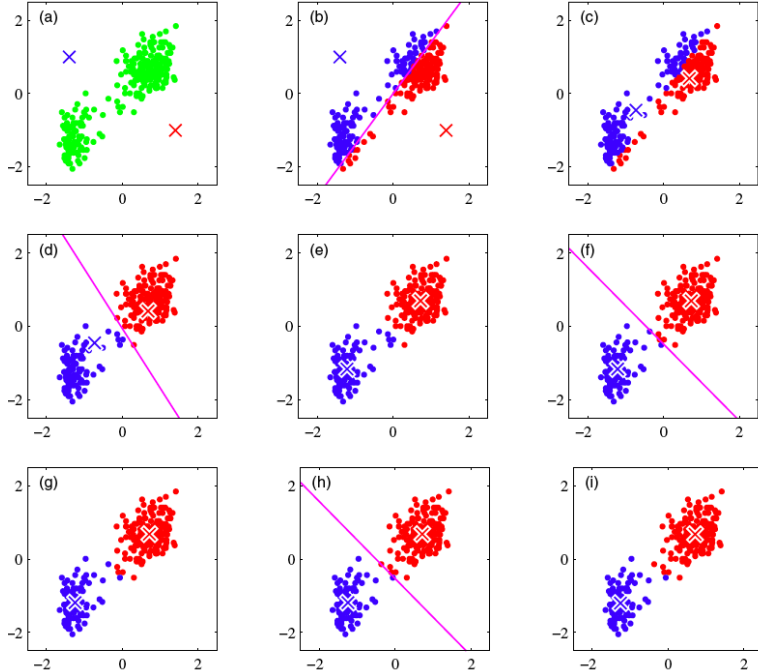



Figure from Bishop

Simple demo: <http://syskall.com/kmeans.js/>


# How to choose K?



[Source]

# Shortcoming of k-means

Convergence of  $k$ -means heavily depends on the **initial** pick of **centroids**. It can perform arbitrarily badly:



## Summary of $K$ -means

We'd love to minimize

$$\sum_{k=1}^K \frac{1}{|C_k|} \sum_{i, i' \in C_k} \|x_i - x_{i'}\|_2^2$$


- It's *infeasible* to actually optimize this in practice, but  $K$ -means at least gives us a so-called **local optimum** of this objective
- The result we get depends both on  $K$ , and also on the *random initialization* that we wind up with
- It's a good idea to **try different random starts** and pick the best result among them
- There's a method called  **$K$ -means++** that improves how the clusters are initialized

# The BFR Algorithm


Extension of  $k$ -means to large data



# BFR Algorithm



- **BFR** [Bradley-Fayyad-Reina] is a variant of  $k$ -means designed to handle **very large** (disk-resident) data sets
- **Assumes** that clusters are normally distributed around a centroid in a Euclidean space
  - Standard deviations in different dimensions may vary
    - Clusters are axis-aligned ellipses
- Goal is to find cluster centroids; point assignment can be done in a second pass through the data.



# BFR Overview

- **Efficient way to summarize clusters:** Want memory required  $O(\text{clusters})$  and not  $O(\text{data})$
- **IDEA: Rather than keeping points, BFR keeps summary statistics of groups of points**
  - 3 sets: Discard set, Compressed set, Retained set
- **Overview of the algorithm:**
  - 1. Initialize  $K$  clusters/centroids
  - 2. Load in a bag of points from disk
  - 3. Assign new points to one of the  $K$  original clusters, if they are within some distance threshold of the cluster
  - 4. Cluster the remaining points, and create new clusters
  - 5. Try to merge new clusters from step 4 with any of the existing clusters
  - 6. Repeat steps 2-5 until all points are examined

# BFR Algorithm

- Points are read from disk one main-memory-full at a time
- Most points from previous memory loads are summarized by simple statistics
- Step 1) From the initial load we select the initial  $k$  centroids by some sensible approach:
  - Take  $k$  random points
  - Take a small random sample and cluster optimally
  - Take a sample; pick a random point, and then  $k-1$  more points, each as far from the previously selected points as possible

# Three Classes of Points

## 3 sets of points which we keep track of:

### ■ Discard set (DS):

- Points close enough to a centroid to be summarized


### ■ Compressed set (CS):

- Groups of points that are close together but not close to any existing centroid
- These points are summarized, but not assigned to a cluster

### ■ Retained set (RS):

- Isolated points waiting to be assigned to a compression set

# BFR: “Galaxies” Picture




**Discard set (DS):** Close enough to a centroid to be summarized  
**Compression set (CS):** Summarized, but not assigned to a cluster  
**Retained set (RS):** Isolated points, we store them as they are

# Summarizing Sets of Points

For each cluster, the discard set (DS) is summarized by:

- The number of points,  **$N$**
- The vector  **$SUM$** , whose  $i^{\text{th}}$  component is the sum of the coordinates of the points in the  $i^{\text{th}}$  dimension
- The vector  **$SUMSQ$** :  $i^{\text{th}}$  component = sum of squares of coordinates in  $i^{\text{th}}$  dimension



A cluster.


All its points are in the **DS**.

The centroid

# Summarizing Points: Comments

- $2d + 1$  values represent any size cluster
  - $d$  = number of dimensions
- Average in **each dimension** (**the centroid**) can be calculated as  $\text{SUM}_i / N$ 
  - $\text{SUM}_i = i^{\text{th}}$  component of SUM
- Variance of a cluster's discard set in dimension  $i$  is:  $(\text{SUMSQ}_i / N) - (\text{SUM}_i / N)^2$ 
  - And standard deviation is the square root of that
- **Next step: Actual clustering**

**Note:** Dropping the “axis-aligned” clusters assumption would require storing full covariance matrix to summarize the cluster. So, instead of **SUMSQ** being a  $d$ -dim vector, it would be a  $d \times d$  matrix, which is too big!



# The “Memory-Load” of Points

## Steps 3-5) Processing “Memory-Load” of points:

- **Step 3)** Find those points that are “**sufficiently close**” to a cluster centroid and add those points to that cluster and the **DS**
  - These points are so close to the centroid that they can be summarized and then discarded
- **Step 4)** Use any in-memory clustering algorithm to cluster the remaining points and the old **RS**
  - Clusters go to the **CS**; outlying points to the **RS**

**Discard set (DS):** Close enough to a centroid to be summarized.

**Compression set (CS):** Summarized, but not assigned to a cluster

**Retained set (RS):** Isolated points



# The “Memory-Load” of Points

## Steps 3-5) Processing “Memory-Load” of points:


- **Step 5) DS set:** Adjust statistics of the clusters to account for the new points
  - Add  $N_s$ ,  $SUM_s$ ,  $SUMSQ_s$
- Consider merging compressed sets in the CS
- **If this is the last round**, merge all compressed sets in the **CS** and all **RS** points into their nearest cluster

**Discard set (DS):** Close enough to a centroid to be summarized.

**Compression set (CS):** Summarized, but not assigned to a cluster

**Retained set (RS):** Isolated points

# Summary: BFR



**Discard set (DS):** Close enough to a centroid to be summarized  
**Compression set (CS):** Summarized, but not assigned to a cluster  
**Retained set (RS):** Isolated points

# A Few Details...

- **Q1)** How do we decide if a point is “close enough” to a cluster that we will add the point to that cluster?
- **Q2)** How do we decide whether two compressed sets (CS) deserve to be combined into one?

# How Close is Close Enough?

- **Q1)** We need a way to decide whether to put a new point into a cluster (and discard)
  - The **Mahalanobis distance** is less than a threshold
  - High likelihood of the point belonging to currently nearest centroid

# Mahalanobis Distance

- **Normalized Euclidean distance from centroid**

- For a given point  $(x_1, \dots, x_d)$  and a given centroid  $(c_1, \dots, c_d)$

1. Normalize in each dimension:  $y_i = (x_i - c_i) / \sigma_i$

2. Take sum of the squares of the  $y_i$

3. Take the square root

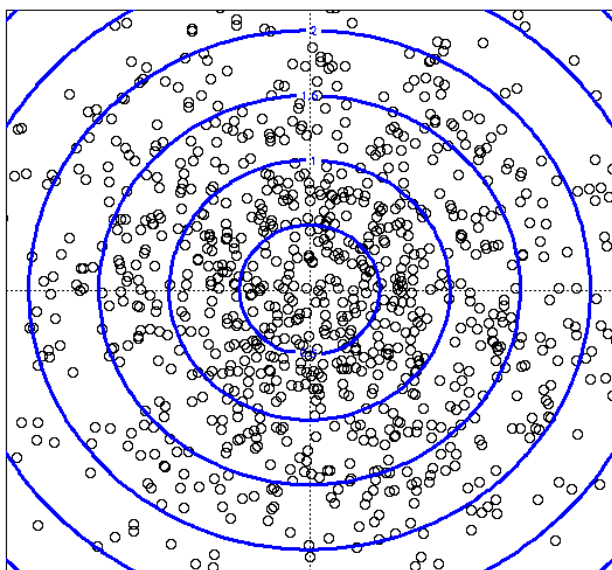
$$d(x, c) = \sqrt{\sum_{i=1}^d \left( \frac{x_i - c_i}{\sigma_i} \right)^2}$$

$\sigma_i$  ... standard deviation of points in the cluster in the  $i^{\text{th}}$  dimension


# Picture: Equal M.D. Regions

## ■ Euclidean vs. Mahalanobis distance


Contours of equidistant points from the origin



Uniformly distributed points,  
Euclidean distance




Normally distributed points,  
Euclidean distance



Normally distributed points,  
Mahalanobis distance

# Picture: Equal M.D. Regions



**Accept a point for a cluster if its M.D. is  $<$  some threshold, e.g., 2 standard deviations**

# Should 2 CS clusters be combined?

## Q2) Should 2 CS clusters be combined?

- Compute the variance of the combined subcluster
  - *N*, *SUM*, and *SUMSQ* allow us to make that calculation quickly
- Combine if the combined variance is below some threshold
- **Many alternatives:** Treat dimensions differently, consider density

